



MATH 4322 Homework 5 Solutions

Cathy Poliak


Fall 2021

Problem 1

The questions relate to the following plots:



- Sketch the tree corresponding to the partition of the predictor space illustrated on the left-hand plot. The numbers inside the boxes indicate the mean of Y within each region.
- Create a diagram similar to the left-hand plot using the tree illustrated in the right-hand plot. You should divide up the predictor space into the correct regions, and indicate the mean for each region.



Problem 2

Suppose we produce ten bootstrapped samples from a data set containing red and green classes. We then apply a classification tree to each bootstrapped sample and, for a specific value of X , produce 10 estimates of $P(\text{Class is Red}|X)$:

0.1, 0.15, 0.2, 0.2, 0.55, 0.6, 0.6, 0.65, 0.7, and 0.75.

There are two common ways to combine these results together into a single class prediction. One is the majority vote approach discussed in this chapter. The second approach is to classify based on the average probability. In this example, what is the final classification under each of these two approaches?

```
px = c(0.1,0.15,0.2,0.2,0.55,0.6,0.6,0.65,0.7,0.75)
#Majority vote
vote = ifelse(px>= 0.5,1,0)
sum(vote)
```

```
## [1] 6
```

```
#Average
mean(px)
```

```
## [1] 0.45
```

With the majority vote we get 6 out of 10 to be Red thus this approach would say that we have Red. The average approach is at 0.45 which is less than 0.5, thus we would say with this approach we have Green.

Problem 3

Provide a detailed explanation of the algorithm that is used to fit a regression tree.

Answer

1. Use recursive binary splitting to grow a large tree on the training data, stopping only when each terminal node has fewer than some minimum number of observations.
2. Apply cost complexity pruning to the large tree in order to obtain a sequence of best subtrees, as a function of α .
3. Use K-fold cross-validation to choose α . That is, divide the training observations into K folds. For each $k = 1, \dots, K$:
 - (a) Repeat Steps 1 and 2 on all but the k th fold of the training data.
 - (b) Evaluate the mean squared prediction error on the data in the left-out k th fold, as a function of α .

Average the results for each value of α , and pick α to minimize the average error.

4. Return the subtree from Step 2 that corresponds to the chosen value of α .

Problem 4

This problem involves the OJ data set which is part of the ISLR package.

- a) Create a training set containing a random sample of 800 observations, and a test set containing the remaining observations.

```
library(ISLR)
data(OJ)
set.seed(1000)
train = sample(nrow(OJ),800)
train.oj = OJ[train,]
test.oj = OJ[-train,]
```

- b) Fit a tree to the training data, with **Purchase** as the response and the other variables as predictors. Use the `summary()` function to produce summary statistics about the tree, and describe the results obtained. What is the training error rate? How many terminal nodes does the tree have?

```
library(tree)
tree.oj = tree(Purchase ~ ., OJ, subset = train)
summary(tree.oj)
```

```
##
## Classification tree:
## tree(formula = Purchase ~ ., data = OJ, subset = train)
## Variables actually used in tree construction:
## [1] "LoyalCH"      "PriceDiff"    "SalePriceMM"
## Number of terminal nodes: 8
## Residual mean deviance: 0.7486 = 592.9 / 792
## Misclassification error rate: 0.16 = 128 / 800
```

Training error rate: 16% Number of terminal nodes: 8

- c) Type in the name of the tree object in order to get a detailed text output. Pick one of the terminal nodes, and interpret the information displayed.


```
tree.oj

## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
## 1) root 800 1066.00 CH ( 0.61500 0.38500 )
##    2) LoyalCH < 0.5036 353 422.60 MM ( 0.28612 0.71388 )
##      4) LoyalCH < 0.276142 170 131.00 MM ( 0.12941 0.87059 )
##        8) LoyalCH < 0.035047 57 10.07 MM ( 0.01754 0.98246 ) *
##        9) LoyalCH > 0.035047 113 108.50 MM ( 0.18584 0.81416 ) *
##      5) LoyalCH > 0.276142 183 250.30 MM ( 0.43169 0.56831 )
##        10) PriceDiff < 0.05 78 79.16 MM ( 0.20513 0.79487 ) *
##        11) PriceDiff > 0.05 105 141.30 CH ( 0.60000 0.40000 ) *
##    3) LoyalCH > 0.5036 447 337.30 CH ( 0.87472 0.12528 )
##      6) LoyalCH < 0.764572 187 206.40 CH ( 0.75936 0.24064 )
##        12) SalePriceMM < 2.125 120 156.60 CH ( 0.64167 0.35833 )
##          24) PriceDiff < -0.35 16 17.99 MM ( 0.25000 0.75000 ) *
##          25) PriceDiff > -0.35 104 126.70 CH ( 0.70192 0.29808 ) *
##        13) SalePriceMM > 2.125 67 17.99 CH ( 0.97015 0.02985 ) *
##      7) LoyalCH > 0.764572 260 91.11 CH ( 0.95769 0.04231 ) *
```

From my node 2): If $LoyalCH < 0.5036$ there are 353 customers with this criteria the deviance is 422.6, the chance that the customer will by Minute Made is 71.388%.

d) Create a plot of the tree, and interpret the results.

```
plot(tree.oj)
text(tree.oj,pretty = 0)
```



The variables that appears to be used to predict if they will buy MM or CH is “LoyalCH”, “SalePriceMM”, and “PriceDiff”.

e) Predict the response on the test data, and produce a confusion matrix comparing the test labels to the predicted test labels. What is the test error rate?

```
tree.pred = predict(tree.oj,test.oj,type = "class")
(con.matrix = table(tree.pred,test.oj$Purchase))
```

```
##
## tree.pred  CH  MM
##          CH 150  38
##          MM  11  71
```

```
#Test error rate
(con.matrix[1,2]+con.matrix[2,1])/sum(con.matrix)
```


```
## [1] 0.1814815
```

- f) Apply the `cv.tree()` function to the training set in order to determine the optimal tree size.
- g) Produce a plot with tree size on the x -axis and cross-validated classification error rate on the y -axis.
- h) Which tree size corresponds to the lowest cross-validated classification error rate?

```
set.seed(2)
cv.oj = cv.tree(tree.oj,FUN = prune.misclass)
cv.oj

## $size
## [1] 8 7 4 2 1
##
## $dev
## [1] 139 139 144 172 308
##
## $k
## [1]      -Inf    0.000000    2.666667   10.500000  151.000000
##
## $method
## [1] "misclass"
##
## attr(,"class")
## [1] "prune"          "tree.sequence"


plot(cv.oj$size,cv.oj$dev,type = "b")
```



It appears that the optimal tree size would be 4.

- i) Produce a pruned tree corresponding to the optimal tree size obtained using cross-validation. If cross-validation does not lead to selection of a pruned tree, then create a pruned tree with five terminal nodes.

```
prune.oj = prune.misclass(tree.oj,best = 4)
plot(prune.oj)
text(prune.oj,pretty = 0)
```



- j) Compare the training error rates between the pruned and unpruned trees. Which is higher?

```
train.tree = predict(tree.oj,type = "class")
(train.matrix = table(train.tree,train.oj$Purchase))
```

```
##
## train.tree  CH  MM
##           CH 450  86
##           MM  42 222
```

```
#Train error rate Un-pruned
(train.matrix[1,2]+train.matrix[2,1])/sum(train.matrix)
```

```
## [1] 0.16
```

```
train.prune = predict(prune.oj,type = "class")
(prune.matrix = table(train.prune,train.oj$Purchase))
```

```
##
## train.prune  CH  MM
##           CH 454  98
##           MM  38 210
```

```
#Train error rate Pruned
(prune.matrix[1,2]+prune.matrix[2,1])/sum(prune.matrix)
```

```
## [1] 0.17
```

The test error rate is higher for the pruned trees.

k) Compare the test error rates between the pruned and unpruned trees. Which is higher?

```
test.tree = predict(tree.oj,test.oj,type = "class")
(test.matrix = table(test.tree,test.oj$Purchase))
```

```
##
## test.tree  CH  MM
##          CH 150  38
##          MM  11  71
```

```
#Train error rate Un-pruned
(test.matrix[1,2]+test.matrix[2,1])/sum(test.matrix)
```

```
## [1] 0.1814815
```

```
test.prune = predict(prune.oj,test.oj,type = "class")
(prune.test.matrix = table(test.prune,test.oj$Purchase))
```

```
##
## test.prune  CH  MM
##          CH 150  44
##          MM  11  65
```

```
#Train error rate Pruned
(prune.test.matrix[1,2]+prune.test.matrix[2,1])/sum(prune.test.matrix)
```

```
## [1] 0.2037037
```

Problem 5

We will use the `Carseats` data set that is in the `ISLR` package to see to predict `Sales` using regression trees and related approaches.

- a) Split the data set into a training set and a test set.


```
set.seed(20)
index = sample(nrow(Carseats), round(0.7*nrow(Carseats)))
train = Carseats[index,]
test = Carseats[-index,]
```

- b) Fit a regression tree to the training set. Plot the tree, and interpret the results. What test MSE do you obtain?

```
tree.carseats = tree(Sales ~ ., train)
summary(tree.carseats)
```

```
##
## Regression tree:
## tree(formula = Sales ~ ., data = train)
## Variables actually used in tree construction:
## [1] "ShelveLoc" "Price" "CompPrice" "Age" "Income"
## [6] "Advertising"
## Number of terminal nodes: 20
## Residual mean deviance: 2.363 = 614.4 / 260
## Distribution of residuals:
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## -4.18400 -0.88550 -0.08422 0.00000 0.95770 4.61000
```

```
plot(tree.carseats)
text(tree.carseats, pretty = 0)
```


```
yhat = predict(tree.carseats,newdata = test)
#Test MSE
(test.mse = mean((yhat - test$Sales)^2))
```

```
## [1] 5.157823
```

There are 20 nodes to this tree. The variables that are used is ShelfLoc, Price, CompPrice, Age, Income and Advertising. With 20 nodes this is very hard to interpret.

- c) Use cross-validation in order to determine the optimal level of tree complexity. Does pruning the tree improve the test MSE?

```
cv.carseats = cv.tree(tree.carseats)
plot(cv.carseats$size, cv.carseats$dev, type = "b")
```



```
cv.carseats
```

```
## $size
## [1] 20 19 18 17 16 15 14 13 11 10 9 8 7 6 5 4 3 2 1
##
## $dev
## [1] 1302.356 1322.618 1322.618 1323.071 1318.090 1318.090 1338.809 1338.809
## [9] 1338.809 1348.762 1343.454 1358.499 1323.423 1350.313 1357.446 1440.282
## [17] 1510.697 1794.510 2316.361
##
## $k
## [1] -Inf 24.80625 24.89140 25.46601 26.76843 26.80908 35.20185
## [8] 35.24160 35.43344 41.76738 43.47571 46.30820 56.97002 72.49995
## [15] 92.24290 110.17098 135.94069 279.25998 531.52217
##
## $method
## [1] "deviance"
##
## attr(,"class")
## [1] "prune" "tree.sequence"
```

It appears that pruning to 7 would be best.

```
prune.carseats = prune.tree(tree.carseats,best = 7)
prune.yhat = predict(prune.carseats,newdata = test)
```

```
#Test MSE
(mse.prune = mean((prune.yhat - test$Sales)^2))
```

```
## [1] 4.679617
```

This does improve the test MSE.

- d) Use the bagging approach in order to analyze this data. What test MSE do you obtain? Use the `importance()` function to determine which variables are most important.

```
library(randomForest)
```

```
## randomForest 4.6-14
```


```
## Type rfNews() to see new features/changes/bug fixes.
```

```
bag.carseat = randomForest(Sales ~., train, mtry = 10, importance = TRUE)
bag.yhat = predict(bag.carseat, newdata = test)
#Test MSE
(bag.mse = mean((bag.yhat - test$Sales)^2))
```

```
## [1] 2.587705
```

```
#Important Variables
varImpPlot(bag.carseat)
```

bag.carseat




The two most important variables are `ShelveLoc` and `Price`.

- (e) Use random forests to analyze this data. What test MSE do you obtain? Use the `importance()` function to determine which variables are most important. Describe the effect of m , the number of variables considered at each split, on the error rate obtained.

```
rf.carseat = randomForest(Sales ~., train, mtry = sqrt(10), importance = TRUE)
rf.yhat = predict(rf.carseat, newdata = test)
#Test MSE
(rf.mse = mean((rf.yhat - test$Sales)^2))
```

```
## [1] 2.89114
```

```
#Important Variables
varImpPlot(rf.carseat)
```



For my random samples, the random forests did not yield much of an improvement over the bagging.

Problem 6

We will use boosting to predict `Salary` in the `Hitters` data set.

- a) Remove the observations for whom the salary information is unknown, and then log-transform the salaries.

```
Hitters2 = na.omit(Hitters)
Hitters2$Salary = log(Hitters2$Salary)
dim(Hitters2)
```

```
## [1] 263 20
```

- b) Create a training set consisting of the first 200 observations, and a test set consisting of the remaining observations.

```
train.hitters = Hitters2[1:200,]
test.hitters = Hitters2[201:263,]
```

- c) Perform boosting on the training set with 1,000 trees. What is the test set MSE? Compare this to the MSE for the regression tree we did in class.

```
library(gbm)
```

```
## Loaded gbm 2.1.8
```


```
set.seed(2)
boost.hit = gbm(Salary ~., train.hitters, distribution = "gaussian",
                n.trees = 1000)
yhat.boost = predict(boost.hit, newdata = test.hitters, n.trees = 1000)
mean((yhat.boost - test.hitters$Salary)^2)
```

```
## [1] 0.2678773
```

This is smaller than the test MSE we received in class.

- d) Which variables appear to be the most important predictors in the boosted model?

```
summary(boost.hit)
```



```
##           var    rel.inf
## CAtBat      CAtBat 17.9553887
## CRuns       CRuns 11.2492102
## CRBI        CRBI  9.6408476
## Walks       Walks  6.9820679
## PutOuts     PutOuts 6.7054644
## CWalks      CWalks  6.4662858
## CHmRun      CHmRun  6.1405886
## Years       Years  5.7824775
## Assists     Assists 5.7148646
## AtBat       AtBat  4.5825777
## RBI         RBI    4.0114290
## Hits        Hits   3.7240791
## HmRun       HmRun   3.1885899
## Runs        Runs   2.6886674
## Errors      Errors  2.3953373
## CHits       CHits   1.4318054
## NewLeague   NewLeague 0.4953537
## League      League  0.4337771
## Division    Division 0.4111881
```

Most important 3 variables: CAtBat, CRuns, and CRBI

e) Now apply bagging to the training set. What is the test set MSE for this approach?

```
bag.hit = randomForest(Salary ~ ., train.hitters, mtry = 19, importance = TRUE)
yhat.bag.hit = predict(bag.hit, test.hitters)
mean((yhat.bag.hit - test.hitters$Salary)^2 )
```

```
## [1] 0.2307637
```

This is a little bit smaller but not much of an improvement over the boosting.