

Lab 17 MATH 4322

Bagging, Random Forest and Boosting

11/09/2021

- We will apply bagging, random forests and boosting to the `Boston` data, using the `randomForest` package.

Question 1: For any data that has p predictors **bagging** requires that we consider how many predictors at each split in a tree?

- `mtry = p`

First, we call the data and create training/testing sets.

```
library(ISLR2)
set.seed(1)
train = sample(1:nrow(Boston), nrow(Boston)/2)
boston.test = Boston[-train, "medv"]
```

Bagging

We perform bagging as follows:

```
library(randomForest)
set.seed(10)
bag.boston = randomForest(medv~., data = Boston,
                           subset = train,
                           mtry = ncol(Boston) - 1, #how many variables
                           importance = TRUE) #what variables are going to be important
bag.boston
```

```
##
## Call:
## randomForest(formula = medv ~ ., data = Boston, mtry = ncol(Boston) - 1, importance = TRUE, su
##               Type of random forest: regression
##               Number of trees: 500
## No. of variables tried at each split: 12
##
##               Mean of squared residuals: 11.5691
##               % Var explained: 84.95
```

Question 2: What is the *MSE* based on the training set?

- $MSE = 11.22857$

How well does this bagged model perform on the test set?


- $\sqrt{11.22857} = 3.350$, which means we are off by \$3.35 thousands dollar.

Question 3: What is the formula to determine the MSE ?

- $MSE = \text{mean}(\text{predicted } y - \text{observed } y)^2$

Run the following in R.

```
yhat.bag = predict(bag.boston,newdata = Boston[-train,])
plot(yhat.bag,boston.test)
abline(0,1)
```



```
mean((yhat.bag - boston.test)^2)
```

```
## [1] 23.23877
```

Question 4: What is the MSE of the test data set?

- $MSE = 23.56$ or $\sqrt{23.56} = \$4.85$ thousand dollars.

We could change the number of trees grown by `randomForest()` using the `ntree` argument:

```
bag.boston = randomForest(medv ~ ., data = Boston,
                           subset = train,
                           mtry = ncol(Boston) - 1,
                           ntree = 25)
```

```
bag.boston
```

```
##
```

```
## Call:
```

```
## randomForest(formula = medv ~ ., data = Boston, mtry = ncol(Boston) - 1, ntree = 25, subset =
```

```
##           Type of random forest: regression
```

```
##           Number of trees: 25
```

```
## No. of variables tried at each split: 12
```

```
##
```

```
##           Mean of squared residuals: 12.30361
```

```
##           % Var explained: 83.99
```

```
yhat.bag = predict(bag.boston, newdata = Boston[-train,])
```

```
mean((yhat.bag - boston.test)^2)
```

```
## [1] 23.06258
```

- The MSE is a little bit higher.

Question 5: What method do we use to get the different trees?

- The bootstrap method

Random Forests

Question 6: For a building a random forest of regression trees, what should be `mtry` (number of predictors to consider at each split)?

- For regression trees the $mtry = p/3$
- For classification trees the $mtry = \sqrt{p}$

Type and run the following in R:

```
set.seed(10)
rf.boston = randomForest(medv ~., data = Boston,
                          subset = train,
                          mtry = (ncol(Boston)-1)/3,
                          importance = TRUE)

yhat.rf = predict(rf.boston, newdata = Boston[-train,])
mean((yhat.rf - boston.test)^2)
```

```
## [1] 18.62328
```

Question 7: Compare the *MSE* of the test data to the *MSE* of the bagging.

* The MSE for the random forest is 19.62 * The MSE for the bagging is 23.56

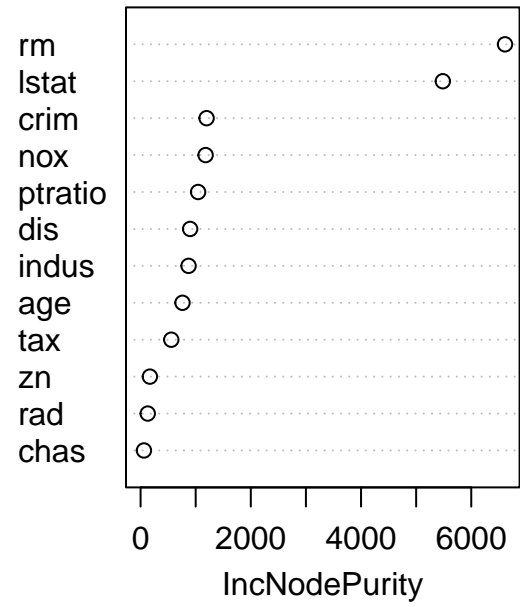
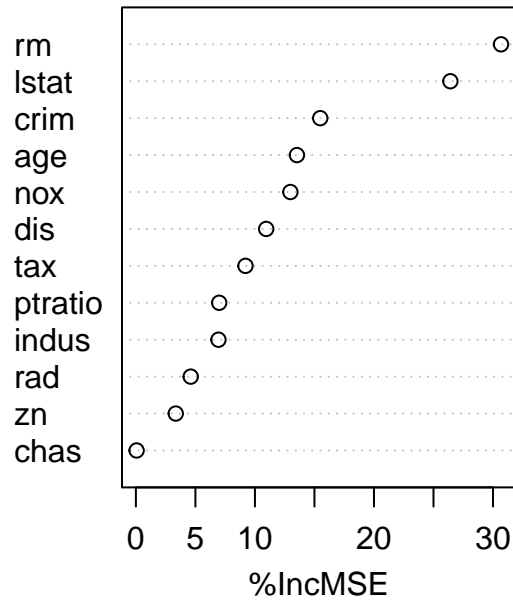
Question 8: Use the `importance()` function what are the two most important variables?

```
importance(rf.boston)
```

```
##           %IncMSE IncNodePurity
## crim    15.48571304    1197.64717
## zn       3.34978057     169.00931
## indus    6.93488857     870.60348
## chas     0.05746934      61.05778
## nox     12.97835448    1179.66670
## rm      30.67206810    6612.55554
## age     13.52685213     760.41982
## dis     10.94707995     899.17273
## rad      4.60598124     129.80949
## tax      9.20624202     556.89248
## ptratio  6.99867017    1044.02812
## lstat    26.41637352    5483.83696
```

```
varImpPlot(rf.boston)
```

rf.boston



- rm and lstat


Boosting

Run the following in R:

```
library(gbm)
```

```
## Loaded gbm 2.1.8
```

```
set.seed(1)
boost.boston = gbm(medv ~., data = Boston[train,],
  distribution = "gaussian", #regression tree, "Bernoulli" - classification
  n.trees = 5000, #default is 500
  interaction.depth = 4) #up to 4 variables that interact with each other
summary(boost.boston)
```



```
##      var      rel.inf
## rm      rm 44.48249588
## lstat    lstat 32.70281223
## crim     crim  4.85109954
## dis      dis  4.48693083
## nox      nox  3.75222394
## age      age  3.19769210
## ptratio  ptratio 2.81354826
```


```
## tax      tax  1.54417603
## indus    indus 1.03384666
## rad      rad  0.87625748
## zn       zn   0.16220479
## chas     chas 0.09671228
```

Question 9: What are the two most important variables with the boosted trees?


- rm and lstat

We can produce *partial dependence plots* for these two variables. The plots illustrate the marginal effect of the selected variables on the response after *integrating* out the other variables.

```
plot(boost.boston, i = "rm")
```



```
plot(boost.boston, i = "lstat")
```



Notice that the house prices are increasing with `rm` and decreasing with `lstat`.

We will use the boosted model to predict `medv` on the test set:

```
yhat.boost = predict(boost.boston,
                      newdata = Boston[-train,],
                      n.trees = 5000)
mean((yhat.boost - boston.test)^2)
```

```
## [1] 18.39057
```

Question 10: Compare this *MSE* to the *MSE* of the random forest and bagging models.

The *MSE* for boosting is lower than random forest and bagging

- The *MSE* for the boosting is 18.84
- The *MSE* for the random forest is 19.62
- The *MSE* for the bagging is 23.56

For the classification we use the confusion matrix proportion of wrong predictor, for the regression we use *MSE* to predict how far we are off by