

Introduction to Neural Networks

Neural Networks


Cathy Poliak, Ph.D.
cpoliak@central.uh.edu

Department of Mathematics
University of Houston

How do neural netWORKs WORK?

Similar to the biological neuron structure, ANNs define the neuron as a:

"Central processing unit that performs a mathematical operation to generate output from a set of inputs."



Essentially, ANN is a set of mathematical function approximations.

Another Example in R

- Data collected at a restaurant through customer interviews. The data set is named `RestuarntTips`
 - The customers were asked to give a score to the following aspects: Service, Ambience, and Food.
 - Also were asked whether they would leave the tip on the basis of these scores, `CustomerWillTip` (Tip = 1 and No-tip = 0)
1. What type of problem is this?
- a) Classification problem
 - b) Regression problem

R Code

Type and run the following in R

```
library(nnet)
#Import dataset RestaurantTips
attach(RestaurantTips)
names(RestaurantTips)

#Train the model based on output from input

model = nnet(CustomerWillTip ~ Service + Ambience + Food,
              data = RestaurantTips,
              size = 5,
              rang = 0.1,
              decay = 5e-2,
              maxit = 5000)

print(model)
```

This output processes the forward and backpropagation until convergence.

The Parameters in the `nnet()` Function

The parameters used in the `nnet()` function can be tuned to improve performance.

- Size : Number of units in the hidden layer
- Decay : Weight decay specifies regularization in the neural network. As a rule of thumb, the more training examples you have, the weaker this term should be. The more parameters you have the higher this term should be.
- Maxit : Maximum iterations, the default is 100.

rang = Initial random weights default is 0.7

Lab Questions

2. How many weights are there?

a) 5

b) 3

c) 15

d) 26

$$(3 \times 5 + 5) + 6$$

3. Type and run the following

```
pred_nnet<-predict(model,RestaurantTips,type = "class")  
(mtab<-table(RestaurantTips$CustomerWillTip,pred_nnet))
```

What is the training error rate?

a) 16.67%

b) 50%

c) 86.67%

d) 4%

```
pred_nnet  
 0 1  
0 12 3  
1 2 13
```

$$\frac{5}{30}$$

Plots of the Model

4. Type and run the following `plotnet(model)`. How many nodes are in the hidden layer?

a) 3

b) 5

c) 1

d) 26

5. Type and run the following `garson(model)`, this is using the `NeuralNetTools`. Which input parameter has the greatest influence to give a tip?

a) Ambience

b) Food

c) Service

d) All of them are equal.

Step-by-step ANN training.

Let us analyze in detail, step by step, all the operations to be done for network training:

1. Initialize the weights w^0 and biases b^0 with random values (one time).
2. Repeat the steps 3 to 5 for each training pattern (presented in random order), until the **error is minimized** or a **stopping criteria is reached**.
3. Conduct **forward propagation** for ANN with weights w^0 and biases b^0 :

input layer \rightarrow hidden layer(s) \rightarrow output layer

4. Conduct **backpropagation**:

compute error \rightarrow take derivative \rightarrow adjust weights/biases

5. Set w^0 and b^0 equal to **adjusted weights/biases**, back to step 3.

The complete pass back and forth is called a **training cycle** or **epoch**. The updated weights and biases are used in the next cycle. We keep recursively training until the error is very minimal.


Training ANN: Minimizing the Error

We desire to adjust the weight and bias parameters such that to **minimize** the error $\hat{y} - y = \text{predicted output} - \text{true output}$.

Example **Single Neuron Model**:

- p nodes x_1, \dots, x_p in the input layer,
- one output node y , calculated as
$$\hat{y} = b + \sum_{j=1}^p w_j x_j$$

Mathematical model



Presume we are supplied with:

- n data samples $\mathbf{x}_i = (x_{1,i}, \dots, x_{p,i}), i = 1, \dots, n$
- n corresponding true responses (or **labels**) $y_i, i = 1, \dots, n$

Training ANN: Minimizing the Error

Feed the **whole data** $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ to our ANN **at once**. Then the **total error** formula is

$$\sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n [y_i - (b + \sum_{j=1}^p w_j x_{j,i})]^2 \equiv f_{err}(b, w_1, \dots, w_p)$$

and we need to **minimize** it with respect to b, w_1, \dots, w_p

Training ANN: Minimizing the Error

Feed the **whole data** $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ to our ANN **at once**. Then the **total error** formula is

$$\sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n [y_i - (b + \sum_{j=1}^p w_j x_{j,i})]^2 \equiv f_{err}(b, w_1, \dots, w_p)$$

and we need to **minimize** it with respect to b, w_1, \dots, w_p

$$\min_{b, w_1, \dots, w_p} \sum_{i=1}^n [y_i - (b + \sum_{j=1}^p w_j x_{j,i})]^2 = \min_{b, w_1, \dots, w_p} f_{err}(b, w_1, \dots, w_p) \quad (1)$$

Similar to the **least squares regression**:

- $\hat{y}_i = \beta_0 + \sum_{j=1}^p \beta_j x_{j,i}$
- $\min_{\beta_0, \dots, \beta_p} \sum_{i=1}^n [y_i - (\beta_0 + \sum_{j=1}^p \beta_j x_{j,i})]^2 = \min_{\beta_0, \dots, \beta_p} f_{err}(\beta_0, \beta_1, \dots, \beta_p)$

which we could **solve analytically** via $\Delta_{\beta=(\beta_0, \dots, \beta_p)} f(\beta) \equiv 0$.

Training ANN: Minimizing the Error

Likewise for the **Linear Neuron Model**, in order to obtain optimal values b, w_1, \dots, w_p that minimize (1), we **could analytically solve**:


$$\Delta_{\mathbf{w}=(b,w_1,\dots,w_p)} f(\mathbf{w}) \equiv 0$$

Training ANN: Minimizing the Error

Likewise for the **Linear Neuron Model**, in order to obtain optimal values b, w_1, \dots, w_p that minimize (1), we **could analytically solve**:

$$\Delta_{\mathbf{w}=(b,w_1,\dots,w_p)} f(\mathbf{w}) \equiv 0$$

$$\Rightarrow \hat{\mathbf{w}} = (\hat{b}, \hat{w}_1, \dots, \hat{w}_p) = \underset{b, w_1, \dots, w_p}{\operatorname{argmin}} \sum_{i=1}^n [y_i - (b + \sum_{j=1}^p w_j x_{j,i})]^2$$



Training ANN: Gradient Descent

While such simple case as **Linear Neuron Model** could be solved **analytically** (read "one could calculate an **explicit formula** for weight and bias estimates"), there are **multiple aspects** to note:

1. We'd like to **avoid solving large systems of equations analytically**.
2. We want a method that **real neurons are likely using**. Hint: they're probably not analytically solving any equations.

Training ANN: Gradient Descent

While such simple case as **Linear Neuron Model** could be solved **analytically** (read "one could calculate an **explicit formula** for weight and bias estimates"), there are **multiple aspects** to note:

1. We'd like to **avoid solving large systems of equations analytically**.
2. We want a method that **real neurons are likely using**. Hint: they're probably not analytically solving any equations.
3. We want a method that **can be generalized** to **multi-layer, non-linear** neural networks, for most of which the closed-form analytical solutions (read "explicit formulas") simply **won't be available**.


That's where such numerical optimization technique as **gradient descent** comes in extremely handy.

Gradient Descent.

Gradient descent is an optimization approach, which entails using:

- Partial derivatives (gradients) of a function, and
- a step size parameter,


in order to calculate the minimum (or maximum) of that function.



We move by the direction of steepest descent.

Gradient Descent for ANN

For neural networks, the gradient descent approach is used when **iterating** the **updates** of weights and biases.



Global cost function in our case is the **squared prediction error**,

$$\frac{1}{2}(\hat{y} - y)^2, \quad \text{or} \quad \frac{1}{2} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

which we are trying to **minimize**.

Gradient Descent: online- and batch-learning

Questions about weight updates:

- Which function to optimize, $\frac{1}{2}(\hat{y} - y)^2$? Or is it $\frac{1}{2} \sum_{i=1}^n (\hat{y}_i - y_i)^2$?
- How often do we need to update the parameter estimates?

Gradient Descent: online- and batch-learning

Questions about weight updates:

- Which function to optimize, $\frac{1}{2}(\hat{y} - y)^2$? Or is it $\frac{1}{2} \sum_{i=1}^n (\hat{y}_i - y_i)^2$?
- How often do we need to update the parameter estimates?

Given data $(\mathbf{x}^1, y^1), \dots, (\mathbf{x}^n, y^n)$, where $\mathbf{x}^i = (x_1^i, \dots, x_p^i)$, gradient descent can be performed for

- **online**-learning - ANN is fed **one training sample (\mathbf{x}^i, y^i) at a time**
 \implies weights are updated each time by minimizing $\frac{1}{2}(\hat{y}^i - y^i)^2$.

Gradient Descent: online- and batch-learning

Questions about weight updates:

- Which function to optimize, $\frac{1}{2}(\hat{y} - y)^2$? Or is it $\frac{1}{2} \sum_{i=1}^n (\hat{y}_i - y_i)^2$?
- How **often** do we need to **update** the parameter estimates?

Given data $(\mathbf{x}^1, y^1), \dots, (\mathbf{x}^n, y^n)$, where $\mathbf{x}^i = (x_1^i, \dots, x_p^i)$, gradient descent can be performed for

- **online**-learning - ANN is fed **one training sample (\mathbf{x}^i, y^i) at a time** \implies weights are updated each time by minimizing $\frac{1}{2}(\hat{y}^i - y^i)^2$.
- **full-batch** - ANN is fed **full training set all at once** \implies weights are updated once **all \hat{y}^i got calculated**, minimize $\frac{1}{2} \sum_{i=1}^n (\hat{y}^i - y^i)^2$.

Gradient Descent: online- and batch-learning

Questions about weight updates:

- Which function to optimize, $\frac{1}{2}(\hat{y} - y)^2$? Or is it $\frac{1}{2} \sum_{i=1}^n (\hat{y}_i - y_i)^2$?
- How **often** do we need to **update** the parameter estimates?

Given data $(\mathbf{x}^1, y^1), \dots, (\mathbf{x}^n, y^n)$, where $\mathbf{x}^i = (x_1^i, \dots, x_p^i)$, gradient descent can be performed for

- **online**-learning - ANN is fed **one training sample (\mathbf{x}^i, y^i) at a time**
 \implies weights are updated each time by minimizing $\frac{1}{2}(\hat{y}^i - y^i)^2$.
- **full-batch** - ANN is fed **full training set all at once** \implies weights are updated once **all \hat{y}^i got calculated**, minimize $\frac{1}{2} \sum_{i=1}^n (\hat{y}^i - y^i)^2$.
- **mini-batches** - ANN is **iteratively** fed **subsets** of training data
 \implies weights are updated once all $\hat{y}^i \in \text{subset}$ got calculated, minimize $\frac{1}{2} \sum_{i \in \text{subset}} (\hat{y}^i - y^i)^2$

Gradient Descent: online- and batch-learning

For either of three approaches (online, full- or mini-batch),

1. The weights & biases are **randomly initialized** at first.

Gradient Descent: online- and batch-learning

For either of three approaches (online, full- or mini-batch),

1. The weights & biases are **randomly initialized** at first.
2. They get updated each time after:
 - a) Data is fed to ANN (be it single sample, full- or mini-batch),
 - b) Forward propagation is performed to get \hat{y}_i 's,
 - c) Error function is calculated (be it single $\frac{1}{2}(\hat{y} - y)$ or $\frac{1}{2} \sum_i (\hat{y}_i - y_i)^2$)
 - d) **Weights are updated via gradient descent.**

Gradient Descent: online- and batch-learning

For either of three approaches (online, full- or mini-batch),

1. The weights & biases are **randomly initialized** at first.
2. They get updated each time after:
 - a) Data is fed to ANN (be it single sample, full- or mini-batch),
 - b) Forward propagation is performed to get \hat{y}_i 's,
 - c) Error function is calculated (be it single $\frac{1}{2}(\hat{y} - y)$ or $\frac{1}{2} \sum_i (\hat{y}_i - y_i)^2$)
 - d) **Weights are updated via gradient descent.**
3. The step 2 is iterated until **all training observations were used**, with **updated weights** being used for **next iteration**.



Gradient Descent: online- and batch-learning

For either of three approaches (online, full- or mini-batch),

1. The weights & biases are **randomly initialized** at first.
2. They get updated each time after:
 - a) Data is fed to ANN (be it single sample, full- or mini-batch),
 - b) Forward propagation is performed to get \hat{y}_i 's,
 - c) Error function is calculated (be it single $\frac{1}{2}(\hat{y} - y)$ or $\frac{1}{2} \sum_i (\hat{y}_i - y_i)^2$)
 - d) **Weights are updated via gradient descent.**
3. The step 2 is iterated until **all training observations were used**, with **updated weights** being used for **next iteration**.
4. Steps 2 & 3 formulate an **epoch** - a **single pass** through the **whole training set**. ANNs are trained over many epochs, with each epoch potentially containing **multiple iterations** (bar the full-batch approach).

Gradient Descent: full- vs and mini-batch example

Example. Below is an exemplary progression for a **full-batch gradient descent** (left) as opposed to **mini-batch** (right). Cost here is the error function, $\frac{1}{2} \sum_i (\hat{y}_i - y_i)$.



Changes of error function are

- fewer and smoother for full-batch,
- more frequent and bumpier for mini-batch.

Iterative Approach: Example

An example to illustrate the iterative method of **online** gradient descent:

- Each day you get lunch at the cafeteria.
 - ▶ Your diet consists of fish, chips, and ketchup.
 - ▶ You get several **portions** (**inputs x**) of each $\Rightarrow x_{fish}, x_{chips}, x_{ketchup}$.

Iterative Approach: Example

An example to illustrate the iterative method of **online** gradient descent:

- Each day you get lunch at the cafeteria.
 - ▶ Your diet consists of fish, chips, and ketchup.
 - ▶ You get several **portions** (**inputs x**) of each $\Rightarrow x_{fish}, x_{chips}, x_{ketchup}$.
- Cashier only tells you the total price of the meal (**response y**).
 - ▶ After several days, you should be able to figure out the **price** (**weight w**) of each portion $\Rightarrow w_{fish}, w_{chips}, w_{ketchup}$


Iterative Approach: Example

An example to illustrate the iterative method of **online** gradient descent:

- Each day you get lunch at the cafeteria.
 - ▶ Your diet consists of fish, chips, and ketchup.
 - ▶ You get several **portions** (**inputs x**) of each $\Rightarrow x_{fish}, x_{chips}, x_{ketchup}$.
- Cashier only tells you the total price of the meal (**response y**).
 - ▶ After several days, you should be able to figure out the **price** (**weight w**) of each portion $\Rightarrow w_{fish}, w_{chips}, w_{ketchup}$
- The iterative approach:
 1. **Start** with **random guesses** for the prices (\Leftrightarrow weights), and then
 2. **Adjust** them to get a better fit to the **observed prices of whole meals** (y_i 's).

Iterative Approach: Example


The **true** weights used by the cashier



$$2(150) + 5(50) + 3(100) = 850$$

Iterative Approach: Example

In reality? True weights are **unknown**.



Iterative Approach: Example

- The portion prices are **weights** of our **linear neuron**.

$$\mathbf{w} = (w_{fish}, w_{chips}, w_{ketchup})$$

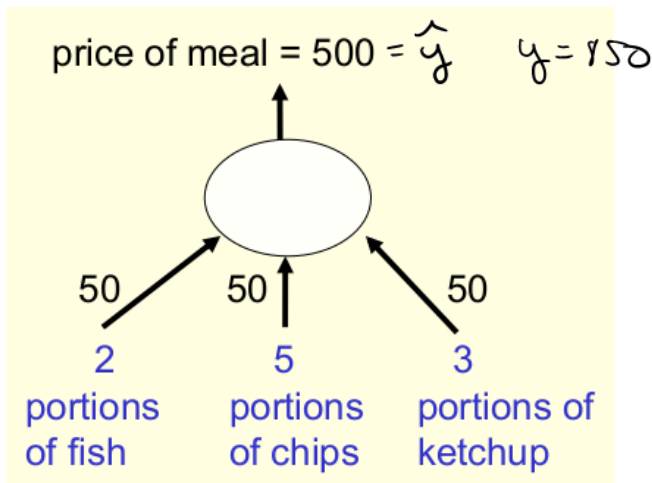
- Each **meal price** is our **response y** , and it is calculated via linear combination of **portion prices** (weights w) multiplied by **# of portions** (inputs x):

$$price = x_{fish}w_{fish} + x_{chips}w_{chips} + x_{ketchup}w_{ketchup}$$

- We start with **random guesses** for the **weights**, and then **adjust the guesses slightly** to give a **better fit** to the prices given by the cashier.

Iterative Approach: Example

Initializing the weights with **random values**, e.g 50, 50, 50.



Iterative Approach: Example

- True value $y = 850$; predicted value $\hat{y} = 500 \implies$

$$\text{Residual error} = 850 - 500 = 350$$

Iterative Approach: Example

- True value $y = 850$; predicted value $\hat{y} = 500 \implies$

$$\text{Residual error} = 850 - 500 = 350$$

- Below we provide a **delta-rule update** for gradient descent algorithm given a **learning rate** α :

$$\Delta w_i = \alpha \times \frac{\partial}{\partial w_i} \left[\frac{1}{2} (\hat{y} - y)^2 \right], \quad i = 1, 2, 3$$

$$\frac{\partial}{\partial w_1} \left[\frac{1}{2} (y - (x_1 w_1 + x_2 w_2 + x_3 w_3))^2 \right]$$

$$= -x_1 (y - (x_1 w_1 + x_2 w_2 + x_3 w_3))$$

$$= -x_1 (y - \hat{y})$$

Iterative Approach: Example

- True value $y = 850$; predicted value $\hat{y} = 500 \implies$

$$\text{Residual error} = 850 - 500 = 350$$

- Below we provide a **delta-rule update** for gradient descent algorithm given a **learning rate** α :

$$\Delta w_i = \alpha \times \frac{\partial}{\partial w_i} \left[\frac{1}{2}(\hat{y} - y)^2 \right], \quad i = 1, 2, 3$$

where

$$\frac{\partial}{\partial w_i} \left[\frac{1}{2}(y - \hat{y})^2 \right] = \frac{\partial}{\partial w_i} \left[\frac{1}{2}(y - \sum_i w_i x_i)^2 \right] = \dots = -x_i(y - \hat{y})$$

Iterative Approach: Example

- True value $y = 850$; predicted value $\hat{y} = 500 \implies$

$$\text{Residual error} = 850 - 500 = 350$$

- Below we provide a **delta-rule update** for gradient descent algorithm given a **learning rate** α :

$$\Delta w_i = \alpha \times \frac{\partial}{\partial w_i} \left[\frac{1}{2} (\hat{y} - y)^2 \right], \quad i = 1, 2, 3$$

where

$$\frac{\partial}{\partial w_i} \left[\frac{1}{2} (y - \hat{y})^2 \right] = \frac{\partial}{\partial w_i} \left[\frac{1}{2} (y - \sum_i w_i x_i)^2 \right] = \dots = -x_i (y - \hat{y})$$

- Hence, the eventual weight updates are:

$$w_i^{upd} = w_i - \Delta w_i = w_i + \alpha \times x_i (y - \hat{y}), \quad i = 1, 2, 3$$


Learning Rate α .

Learning rate α is a scalar parameter used to set the rate of adjustments/updates in order to reduce the training errors faster.

Picking a learning rate value is an art, which can lead to your model:

- training & learning fast,
- training & learning slow,
- not training & learning at all.

Unfortunately, we won't be mastering that art in this course.



Iterative Approach: Example

- If we select learning rate $\alpha = \frac{1}{35}$:

- ▶ $w_1^{upd} = 50 + \frac{1}{35} 2(350) = 70$,
- ▶ $w_2^{upd} = 50 + \frac{1}{35} 5(350) = 100$,
- ▶ $w_3^{upd} = 50 + \frac{1}{35} 3(350) = 80$

$$w_i^{upd} = w_i + \alpha x_i (y - \hat{y})$$

Iterative Approach: Example

- If we select learning rate $\alpha = \frac{1}{35}$:
 - ▶ $w_1^{upd} = 50 + \frac{1}{35}2(350) = 70$,
 - ▶ $w_2^{upd} = 50 + \frac{1}{35}5(350) = 100$,
 - ▶ $w_3^{upd} = 50 + \frac{1}{35}3(350) = 80$
- The updated predicted value \hat{y} :

$$\hat{y} = \sum_i w_i^{upd} x_i = 880,$$

which is much closer to the true value $y = 850$.

$$2(70) + 5(100) + 3(80) = 880$$

$$\text{residual} = 880 - 850 = -30$$

Iterative Approach: Example

- If we select learning rate $\alpha = \frac{1}{35}$:
 - ▶ $w_1^{upd} = 50 + \frac{1}{35} 2(350) = 70$,
 - ▶ $w_2^{upd} = 50 + \frac{1}{35} 5(350) = 100$,
 - ▶ $w_3^{upd} = 50 + \frac{1}{35} 3(350) = 80$
- The updated predicted value \hat{y} :

$$\hat{y} = \sum_i w_i^{upd} x_i = 880,$$

which is much closer to the true value $y = 850$.

- Notice that, while the weight for **fish** and **ketchup** got **better** (70 is closer to 150, 80 is closer to 100), the weight for **chips** got **worse** (100 instead of correct value 50)!
- We aren't guaranteed a weight improvement at each update, but over time, given sufficient data, weights typically converge to the correct values.

$$\omega_1^{upd^2} = 70 + \frac{1}{35} 2(-30) = 64.28$$

$$\omega_2^{upd^2} = 100 + \frac{1}{35} (5)(-30) = 95.71$$

$$\omega_3^{upd^3} = 80 + \frac{1}{35} (3)(-30) = 77.43$$


Linear Neuron Model: Limitations

Question. If we stick with a **Linear Neuron Model**, where each neuron simply outputs a weighted **linear combination** of its outputs, what type of function of original inputs x_1, \dots, x_n would we inevitably get?

Linear Neuron Model: Limitations

Question. If we stick with a **Linear Neuron Model**, where each neuron simply outputs a weighted **linear combination** of its outputs, what type of function of original inputs x_1, \dots, x_n would we inevitably get?

Answer. **Linear function.**



Linear functions **can only do that much** when dealing with **non-linearity**.